

Variational Inference for Diffusion Models

Tomonari MASADA

masada@rikkyo.ac.jp

Contents

Marginal likelihood

Variational lower bound

Variational posterior

Generative modeling of observations

Appendix

Marginal likelihood

We consider a Bayesian generative model whose joint distribution is

$$p_{\theta}(\mathbf{x}_{0:T}) = p_{\theta}(\mathbf{x}_T) \prod_{t=1}^T p_{\theta}(\mathbf{x}_{t-1} | \mathbf{x}_t) \quad (1)$$

where \mathbf{x}_0 is an observation and $\{\mathbf{x}_t : t = 1, \dots, T\}$ are latent random variables. The distribution of \mathbf{x}_t is only conditioned on \mathbf{x}_{t+1} . The log of marginal likelihood (= evidence) of \mathbf{x}_0 is

$$\log p_{\theta}(\mathbf{x}_0) = \log \int p_{\theta}(\mathbf{x}_{0:T}) d\mathbf{x}_{1:T} \quad (2)$$

Contents

Marginal likelihood

Variational lower bound

Variational posterior

Generative modeling of observations

Appendix

ELBO

Jensen's inequality gives a lower bound of $\log p_\theta(\mathbf{x}_0)$ as follows:

$$\begin{aligned}\log p_\theta(\mathbf{x}_0) &= \log \int q_\psi(\mathbf{x}_{1:T}|\mathbf{x}_0) \frac{p_\theta(\mathbf{x}_{0:T})}{q_\psi(\mathbf{x}_{1:T}|\mathbf{x}_0)} d\mathbf{x}_{1:T} \\ &\geq \int q_\psi(\mathbf{x}_{1:T}|\mathbf{x}_0) \log \frac{p_\theta(\mathbf{x}_{0:T})}{q_\psi(\mathbf{x}_{1:T}|\mathbf{x}_0)} d\mathbf{x}_{1:T} \\ &= \int q_\psi(\mathbf{x}_{1:T}|\mathbf{x}_0) \log \frac{p_\theta(\mathbf{x}_{0:T})}{q_\psi(\mathbf{x}_{1:T}|\mathbf{x}_0)} d\mathbf{x}_{1:T} \equiv L_{\text{VLB}}\end{aligned}\quad (3)$$

where $q_\psi(\mathbf{x}_{1:T}|\mathbf{x}_0)$ is a variational posterior, which is conditioned on the observation \mathbf{x}_0 as in VAE. We train our model over many observations $\mathcal{X} \equiv \{\mathbf{x}_0^{(1)}, \dots, \mathbf{x}_0^{(N)}\}$ in an amortized manner.

Factorization assumption

We assume that $q_\psi(\mathbf{x}_{1:T}|\mathbf{x}_0)$ factorizes as

$$q_\psi(\mathbf{x}_{1:T}|\mathbf{x}_0) = q_\psi(\mathbf{x}_1|\mathbf{x}_0) \prod_{t=2}^T q_\psi(\mathbf{x}_t|\mathbf{x}_{t-1}, \mathbf{x}_0) \quad (4)$$

Then the variational lower bound L_{VLB} can be rewritten as follows:

$$\begin{aligned} L_{\text{VLB}} &= \int q_\psi(\mathbf{x}_{1:T}|\mathbf{x}_0) \log \frac{p_\theta(\mathbf{x}_{0:T})}{q_\psi(\mathbf{x}_{1:T}|\mathbf{x}_0)} d\mathbf{x}_{1:T} \\ &= \int q_\psi(\mathbf{x}_{1:T}|\mathbf{x}_0) \log \frac{p_\theta(\mathbf{x}_T) \prod_{t=1}^T p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q_\psi(\mathbf{x}_1|\mathbf{x}_0) \prod_{t=2}^T q_\psi(\mathbf{x}_t|\mathbf{x}_{t-1}, \mathbf{x}_0)} d\mathbf{x}_{1:T} \\ &= \int q_\psi(\mathbf{x}_{1:T}|\mathbf{x}_0) \log p_\theta(\mathbf{x}_T) d\mathbf{x}_{1:T} + \int q_\psi(\mathbf{x}_{1:T}|\mathbf{x}_0) \sum_{t=2}^T \log \frac{p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q_\psi(\mathbf{x}_t|\mathbf{x}_{t-1}, \mathbf{x}_0)} d\mathbf{x}_{1:T} \\ &\quad + \int q_\psi(\mathbf{x}_{1:T}|\mathbf{x}_0) \log \frac{p_\theta(\mathbf{x}_0|\mathbf{x}_1)}{q_\psi(\mathbf{x}_1|\mathbf{x}_0)} d\mathbf{x}_{1:T} \end{aligned} \quad (5)$$

Using Bayes' rule, we have

$$q_\psi(\mathbf{x}_t|\mathbf{x}_{t-1}, \mathbf{x}_0) = \frac{q_\psi(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)q_\psi(\mathbf{x}_t|\mathbf{x}_0)}{q_\psi(\mathbf{x}_{t-1}|\mathbf{x}_0)} \quad (6)$$

We replace $q_\psi(\mathbf{x}_t|\mathbf{x}_{t-1}, \mathbf{x}_0)$ appearing in L_{VLB} with $\frac{q_\psi(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)q_\psi(\mathbf{x}_t|\mathbf{x}_0)}{q_\psi(\mathbf{x}_{t-1}|\mathbf{x}_0)}$ based on Eq. (6) and rewrite L_{VLB} as follows:

$$\begin{aligned} L_{\text{VLB}} = & \int q_\psi(\mathbf{x}_{1:T}|\mathbf{x}_0) \log p_\theta(\mathbf{x}_T) d\mathbf{x}_{1:T} + \int q_\psi(\mathbf{x}_{1:T}|\mathbf{x}_0) \sum_{t=2}^T \log \frac{p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q_\psi(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)} d\mathbf{x}_{1:T} \\ & + \int q_\psi(\mathbf{x}_{1:T}|\mathbf{x}_0) \sum_{t=2}^T \log \frac{q_\psi(\mathbf{x}_{t-1}|\mathbf{x}_0)}{q_\psi(\mathbf{x}_t|\mathbf{x}_0)} d\mathbf{x}_{1:T} + \int q_\psi(\mathbf{x}_{1:T}|\mathbf{x}_0) \log \frac{p_\theta(\mathbf{x}_0|\mathbf{x}_1)}{q_\psi(\mathbf{x}_1|\mathbf{x}_0)} d\mathbf{x}_{1:T} \quad (7) \end{aligned}$$

(cont.)

$$\begin{aligned}
L_{\text{VLB}} &= \int q_{\psi}(\mathbf{x}_{1:T}|\mathbf{x}_0) \log \frac{p_{\theta}(\mathbf{x}_T)}{q_{\psi}(\mathbf{x}_T|\mathbf{x}_0)} d\mathbf{x}_{1:T} + \sum_{t=2}^T \int q_{\psi}(\mathbf{x}_{1:T}|\mathbf{x}_0) \log \frac{p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q_{\psi}(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)} d\mathbf{x}_{1:T} \\
&\quad + \int q_{\psi}(\mathbf{x}_{1:T}|\mathbf{x}_0) \log p_{\theta}(\mathbf{x}_0|\mathbf{x}_1) d\mathbf{x}_{1:T} \equiv L_T + \sum_{t=2}^T L_{t-1} + L_0
\end{aligned} \tag{8}$$

We can rewrite L_{t-1} as follows:

$$\begin{aligned}
L_{t-1} &\equiv \int q_{\psi}(\mathbf{x}_{1:T}|\mathbf{x}_0) \log \frac{p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q_{\psi}(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)} d\mathbf{x}_{1:T} \\
&= \int q_{\psi}(\mathbf{x}_1|\mathbf{x}_0) \prod_{t' \neq t} q_{\psi}(\mathbf{x}_{t'-1}|\mathbf{x}_{t'}, \mathbf{x}_0) \left(q_{\psi}(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) \log \frac{p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q_{\psi}(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)} \right) d\mathbf{x}_{1:T} \\
&= - \int q_{\psi}(\mathbf{x}_1|\mathbf{x}_0) \prod_{t' \neq t} q_{\psi}(\mathbf{x}_{t'-1}|\mathbf{x}_{t'}, \mathbf{x}_0) D_{\text{KL}}(q_{\psi}(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) \parallel p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t)) d\mathbf{x}_{1:T} \\
&\equiv -\mathbb{E}_{\neg t} [D_{\text{KL}}(q_{\psi}(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) \parallel p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t))]
\end{aligned} \tag{9}$$

Contents

Marginal likelihood

Variational lower bound

Variational posterior

Generative modeling of observations

Appendix

Tractable variational posterior

We define our variational posterior with parameters $\psi \equiv \{\alpha_t : t = 1, \dots, T\}$ as follows:

$$q_\psi(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_t; \sqrt{\alpha_t} \mathbf{x}_{t-1}, (1 - \alpha_t) \mathbf{I}) \quad (10)$$

Therefore, $q_\psi(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{x}_0) = q_\psi(\mathbf{x}_t | \mathbf{x}_{t-1})$ for $t > 1$.

Based on Eq. (24), we can obtain $q_\psi(\mathbf{x}_t | \mathbf{x}_{t-2})$ as follows:

$$\begin{aligned} q_\psi(\mathbf{x}_t | \mathbf{x}_{t-2}) &= \mathcal{N}(\mathbf{x}_t; \sqrt{\alpha_t \alpha_{t-1}} \mathbf{x}_{t-2}, ((1 - \alpha_t) + \alpha_t(1 - \alpha_{t-1})) \mathbf{I}) \\ &= \mathcal{N}(\mathbf{x}_t; \sqrt{\alpha_t \alpha_{t-1}} \mathbf{x}_{t-2}, (1 - \alpha_t \alpha_{t-1}) \mathbf{I}) \end{aligned} \quad (11)$$

By repeating the same argument, we obtain $q_\psi(\mathbf{x}_t | \mathbf{x}_0)$ as follows:

$$q_\psi(\mathbf{x}_t | \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_t; \sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t) \mathbf{I}) \quad (12)$$

where $\bar{\alpha}_t = \prod_{s=1}^t \alpha_s$. It is easy to sample from $q_\psi(\mathbf{x}_t | \mathbf{x}_0)$.

We regard ψ as free parameters and drop ψ from our notations from now on.

We rewrite $q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)$ in L_{t-1} as follows:

$$\begin{aligned}
q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) &= \frac{q(\mathbf{x}_t|\mathbf{x}_{t-1}, \mathbf{x}_0)q(\mathbf{x}_{t-1}|\mathbf{x}_0)}{q(\mathbf{x}_t|\mathbf{x}_0)} \quad (\text{by using Bayes' rule}) \\
&\propto \exp \left(-\frac{1}{2} \left(\frac{(\mathbf{x}_t - \sqrt{\alpha_t} \mathbf{x}_{t-1})^2}{1 - \alpha_t} + \frac{(\mathbf{x}_{t-1} - \sqrt{\bar{\alpha}_{t-1}} \mathbf{x}_0)^2}{1 - \bar{\alpha}_{t-1}} - \frac{(\mathbf{x}_t - \sqrt{\bar{\alpha}_t} \mathbf{x}_0)^2}{1 - \bar{\alpha}_t} \right) \right) \\
&\propto \exp \left(-\frac{1}{2} \left(\left(\frac{\alpha_t}{1 - \alpha_t} + \frac{1}{1 - \bar{\alpha}_{t-1}} \right) \mathbf{x}_{t-1}^2 - 2 \left(\frac{\sqrt{\alpha_t}}{1 - \alpha_t} \mathbf{x}_t + \frac{\sqrt{\bar{\alpha}_{t-1}}}{1 - \bar{\alpha}_{t-1}} \mathbf{x}_0 \right) \mathbf{x}_{t-1} \right) \right)
\end{aligned} \tag{13}$$

We denote the element-wise mean and variance of $q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)$ by $\tilde{\boldsymbol{\mu}}(\mathbf{x}_t, \mathbf{x}_0)$ and $\tilde{\beta}_t$ respectively. That is,

$$q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_{t-1}; \tilde{\boldsymbol{\mu}}(\mathbf{x}_t, \mathbf{x}_0), \tilde{\beta}_t) \tag{14}$$

Then

$$\tilde{\beta}_t = 1 / \left(\frac{\alpha_t}{1 - \alpha_t} + \frac{1}{1 - \bar{\alpha}_{t-1}} \right) = \frac{(1 - \alpha_t)(1 - \bar{\alpha}_{t-1})}{\alpha_t - \alpha_t \bar{\alpha}_{t-1} + 1 - \alpha_t} = \frac{1 - \bar{\alpha}_{t-1}}{1 - \bar{\alpha}_t} (1 - \alpha_t) \tag{15}$$

$$\begin{aligned}
\tilde{\boldsymbol{\mu}}(\mathbf{x}_t, \mathbf{x}_0) &= \left(\frac{\sqrt{\alpha_t}}{1 - \alpha_t} \mathbf{x}_t + \frac{\sqrt{\bar{\alpha}_{t-1}}}{1 - \bar{\alpha}_{t-1}} \mathbf{x}_0 \right) / \left(\frac{\alpha_t}{1 - \alpha_t} + \frac{1}{1 - \bar{\alpha}_{t-1}} \right) \\
&= \left(\frac{\sqrt{\alpha_t}}{1 - \alpha_t} \mathbf{x}_t + \frac{\sqrt{\bar{\alpha}_{t-1}}}{1 - \bar{\alpha}_{t-1}} \mathbf{x}_0 \right) \frac{1 - \bar{\alpha}_{t-1}}{1 - \bar{\alpha}_t} (1 - \alpha_t) \\
&= \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} \mathbf{x}_t + \frac{\sqrt{\bar{\alpha}_{t-1}}(1 - \alpha_t)}{1 - \bar{\alpha}_t} \mathbf{x}_0
\end{aligned} \tag{16}$$

Based on Eq. (12), we reparameterize \mathbf{x}_t as $\mathbf{x}_t(\mathbf{x}_0, \boldsymbol{\epsilon}) = \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + (1 - \bar{\alpha}_t) \boldsymbol{\epsilon}$ for $\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ and rewrite $\tilde{\boldsymbol{\mu}}(\mathbf{x}_t, \mathbf{x}_0)$ as follows:

$$\begin{aligned}
\tilde{\boldsymbol{\mu}}(\mathbf{x}_t, \mathbf{x}_0) &= \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} \mathbf{x}_t + \frac{\sqrt{\bar{\alpha}_{t-1}}(1 - \alpha_t)}{1 - \bar{\alpha}_t} \left(\frac{\mathbf{x}_t}{\sqrt{\bar{\alpha}_t}} - \frac{\sqrt{1 - \bar{\alpha}_t}}{\sqrt{\bar{\alpha}_t}} \boldsymbol{\epsilon} \right) \\
&= \left(\frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} + \frac{1 - \alpha_t}{(1 - \bar{\alpha}_t)\sqrt{\alpha_t}} \right) \mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}\sqrt{\alpha_t}} \boldsymbol{\epsilon} \\
&= \frac{1}{\sqrt{\alpha_t}} \left(\left(\frac{\alpha_t(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} + \frac{1 - \alpha_t}{1 - \bar{\alpha}_t} \right) \mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \boldsymbol{\epsilon} \right) \\
&= \frac{1}{\sqrt{\alpha_t}} \left(\mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \boldsymbol{\epsilon} \right)
\end{aligned} \tag{17}$$

Contents

Marginal likelihood

Variational lower bound

Variational posterior

Generative modeling of observations

Appendix

Generative modeling of observations

Here we specify the details of our Bayesian generative model firstly in this presentation.

$$p_{\theta}(\mathbf{x}_T) = \mathcal{N}(\mathbf{x}_T; \mathbf{0}, \mathbf{I}) \quad (18)$$

$$p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t) = \mathcal{N}(\mathbf{x}_{t-1}; \boldsymbol{\mu}_{\theta}(\mathbf{x}_t, t), \boldsymbol{\Sigma}_{\theta}(\mathbf{x}_t, t)) \quad (19)$$

We assume that $\boldsymbol{\Sigma}_{\theta}(\mathbf{x}_t, t) = \sigma_t^2 \mathbf{I}$ as discussed in [1].

Based on Eqs. (14) and (19), it can be said that the KL divergence appearing in L_{t-1} in Eq. (9) is a KL divergence from one Gaussian distribution to another.

Therefore, we can rewrite L_{t-1} as follows:

$$L_{t-1} = -\mathbb{E}_{\neg t} \left[\frac{1}{2\sigma_t^2} \|\tilde{\boldsymbol{\mu}}(\mathbf{x}_t, \mathbf{x}_0) - \boldsymbol{\mu}_\theta(\mathbf{x}_t, t)\|^2 \right] + \text{const.} \quad (20)$$

By using the reparameterization $\mathbf{x}_t(\mathbf{x}_0, \epsilon) = \sqrt{\bar{\alpha}_t}\mathbf{x}_0 + (1 - \bar{\alpha}_t)\epsilon$ and Eq. (17), we further rewrite L_{t-1} as follows:

$$L_{t-1} = -\mathbb{E}_{\neg t} \left[\frac{1}{2\sigma_t^2} \left\| \frac{1}{\sqrt{\alpha_t}} \left(\mathbf{x}_t(\mathbf{x}_0, \epsilon) - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon \right) - \boldsymbol{\mu}_\theta(\mathbf{x}_t(\mathbf{x}_0, \epsilon), t) \right\|^2 \right] + \text{const.} \quad (21)$$

We may parameterize $\boldsymbol{\mu}_\theta(\mathbf{x}_t, t)$ as follows [1]:

$$\boldsymbol{\mu}_\theta(\mathbf{x}_t, t) = \frac{1}{\sqrt{\alpha_t}} \left(\mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(\mathbf{x}_t, t) \right) \quad (22)$$

where ϵ_θ is a function approximator intended to predict ϵ from \mathbf{x}_t .

By using parameterization, we can rewrite L_{t-1} as follows:

$$L_{t-1} = -\mathbb{E}_{\neg t} \left[\frac{(1 - \alpha_t)^2}{2\sigma_t^2(1 - \bar{\alpha}_t)} \|\boldsymbol{\epsilon} - \boldsymbol{\epsilon}_\theta(\sqrt{\bar{\alpha}_t}\mathbf{x}_0 + (1 - \bar{\alpha}_t)\boldsymbol{\epsilon}, t)\|^2 \right] + \text{const.} \quad (23)$$

Contents

Marginal likelihood

Variational lower bound

Variational posterior

Generative modeling of observations

Appendix

Appendix 1

$$\begin{aligned} & \int \exp \left(-\frac{(x-ay)^2}{2s^2} - \frac{(y-bz)^2}{2t^2} \right) dy = \int \exp \left(-\frac{t^2(x-ay)^2 + s^2(y-bz)^2}{2s^2t^2} \right) dy \\ &= \int \exp \left(-\frac{(s^2 + t^2a^2)y^2 - 2(s^2bz + t^2ax)y + t^2x^2 + s^2b^2z^2}{2s^2t^2} \right) dy \\ &= \exp \left(-\frac{t^2x^2 + s^2b^2z^2}{2s^2t^2} \right) \int \exp \left(-\frac{s^2 + t^2a^2}{2s^2t^2} \left(y^2 - \frac{2(s^2bz + t^2ax)}{s^2 + t^2a^2} y \right) \right) dy \\ &= \exp \left(-\frac{t^2x^2 + s^2b^2z^2}{2s^2t^2} + \frac{(s^2bz + t^2ax)^2}{2s^2t^2(s^2 + t^2a^2)} \right) \int \exp \left(-\frac{s^2 + t^2a^2}{2s^2t^2} \left(y - \frac{s^2bz + t^2ax}{s^2 + t^2a^2} \right)^2 \right) dy \\ &\propto \exp \left(-\frac{s^2t^2x^2 + s^4b^2z^2 + t^4a^2x^2 + s^2t^2a^2b^2z^2 - t^4a^2x^2 - 2s^2t^2abzx - s^4b^2z^2}{2s^2t^2(s^2 + t^2a^2)} \right) \\ &= \exp \left(-\frac{x^2 - 2abzx + a^2b^2z^2}{2(s^2 + t^2a^2)} \right) = \exp \left(-\frac{(x-abz)^2}{2(s^2 + t^2a^2)} \right) \end{aligned} \tag{24}$$



Jonathan Ho, Ajay Jain, and Pieter Abbeel.

Denoising diffusion probabilistic models.

CoRR, [abs/2006.11239](https://arxiv.org/abs/2006.11239), 2020.